

Enterprise Continuous Monitoring and Threat Detection on AWS

Author: Adedayo

Specialization: Cloud Security & Threat Detection

Platform: Amazon Web Services (AWS)

1. Problem Statement.

A SaaS organization required continuous threat detection across its AWS environment. Limited monitoring and delayed alerting increased security risk. A centralized detection framework was required.

I implemented this solution to improve visibility into API activity, network traffic, security misconfigurations, and threat indicators, while enabling rapid incident response.

2. Objectives

The primary objectives were to:

- Enable continuous security monitoring
- Detect suspicious IAM and API activity
- Identify network-based threats
- Centralize security findings
- Automate low-severity remediation
- Improve incident response readiness

3. Security Monitoring Architecture

The monitoring framework was built using the following AWS services:

- Amazon GuardDuty
- AWS CloudTrail
- AWS Security Hub
- VPC Flow Logs
- Amazon CloudWatch
- Amazon SNS
- AWS Lambda

All security events and findings were centralized for analysis and response.

4. GuardDuty Configuration

I enabled Amazon GuardDuty to monitor:

The screenshot shows the AWS GuardDuty Summary page. On the left, there's a sidebar with options like 'Summary', 'Findings', 'Malware scans', 'Protection plans' (with sub-options for S3, EKS, Extended Threat Detection, Runtime Monitoring, Malware Protection, RDS, and Lambda), 'Accounts', 'Usage', 'Suppression rules', 'Settings', and 'Lists'. The main area has a 'Summary' card with the message 'View and analyze security trends based on GuardDuty findings in your AWS environment.' It shows 'Updated a few seconds ago' and a 'Today' dropdown. Below this is an 'Overview' section with four boxes: 'Attack sequences - new' (0), 'Total findings' (0), 'Resources with findings' (0), and 'Accounts with findings' (0). To the right is a 'Findings' section with a grid for Critical, High, Medium, and Low severity levels, all showing 0 findings. A tooltip says 'Prioritize triaging and remediating topmost severity detections.' Below the findings is a callout for 'Introducing the new AWS Security Hub'.

- CloudTrail logs
- VPC Flow Logs
- DNS activity

GuardDuty was configured to generate findings for suspicious behavior and compromised credentials.

An SNS topic was created to distribute GuardDuty alerts to the security team.

The screenshot shows the AWS SNS Topics page. The left sidebar includes 'Dashboard', 'Topics' (selected), and 'Subscriptions'. Under 'Topics', there are sections for 'Mobile' (Push notifications, Text messaging (SMS)) and 'Amazon SNS' (Topic name: GuarddutyFindings, ARN: arnawsnsus-east-1:879381257906:GuarddutyFindings, Type: Standard). A blue banner at the top says 'New Feature: Amazon SNS now supports High Throughput FIFO topics. Learn more.' Below the topic details are tabs for 'Subscriptions', 'Access policy', 'Data protection policy', 'Delivery policy (HTTP/S)', 'Delivery status logging', 'Encryption', and 'Tags'. The 'Subscriptions' tab is selected, showing one entry: 'Subscription ID: 1, Endpoint: arnawsnsus-east-1:879381257906:GuarddutyFindings, Status: Pending confirmation, Protocol:'. There are buttons for 'Edit', 'Delete', 'Request confirmation', 'Confirm subscription', and 'Create subscription'.

Sample findings were generated to validate alert delivery.

5. Automated Remediation for Low-Severity Findings

I implemented automated remediation for low-severity findings using AWS Lambda.

The screenshot shows the AWS Lambda function editor for a function named "LowSeverityGuardDutyRemediation". The code is written in Python and performs the following steps:

```
lambda_function.py
1 import json
2 import boto3
3
4 ec2 = boto3.client('ec2')
5
6 def lambda_handler(event, context):
7     print("Receive event:", json.dumps(event))
8     finding = event['detail']
9     severity = finding['severity']
10
11     # Only process low severity
12     if severity > 4:
13         print("Skipping non low severity finding")
14         return
15
16     resources = finding.get('resources', [])
17     for resource in resources:
18         if resource['type'] == 'AWS::EC2::Instance':
19             instance_id = resource['id']
20             sg_response = ec2.describe_instances(InstanceIds=[instance_id])
21             security_groups = sg_response['Reservations'][0]['Instances'][0]['SecurityGroups']
22
23             ip_address = finding.get('service', {}).get('action', {}).get('remoteIpDetails', {}).get('ipAddress')
24             if ip_address:
25                 for sg in security_groups:
26                     sg_id = sg['GroupId']
27                     try:
28                         ec2.revoke_security_group_ingress(sg_id, ip_address)
```

A dedicated IAM role was created with permissions to modify tagged security groups.

The screenshot shows the AWS IAM Role policy editor for a role named "LambdaLowSeverityRemediationRole". The policy document is defined in JSON:

```
1 Version: "2012-10-17",
2 Statement: [
3     {
4         Effect: "Allow",
5         Action: [
6             "ec2:DescribeInstances",
7             "ec2:DescribeSecurityGroups",
8             "ec2:RevokeSecurityGroupIngress",
9             "ec2:AuthorizeSecurityGroupIngress",
10            "logs:CreateLogGroup",
11            "logs:CreateLogStream",
12            "logs:PutLogEvents"
13        ],
14        Resource: "*"
15    }
16 ]
17 }
```

In production environments, permissions would be restricted to specific ARNs and tagged resources to enforce least privilege.

A Lambda function was deployed and triggered using EventBridge when low-severity findings occurred.

6. CloudTrail and API Monitoring

I configured a centralized CloudTrail trail to capture all AWS API activity.

Logs were forwarded to CloudWatch Logs for near real-time analysis.

The screenshot shows the AWS CloudTrail console with the following details:

- General details:**
 - Trail logging: Enabled
 - Trail name: SecureSphereCloudtrail
 - Multi-region trail: Yes
 - Apply trail to my organization: Not enabled
- CloudWatch Logs:**
 - Log group: aws-cloudtrail-logs-securesphere
- IAM Role:** arn:aws:iam::879581257906:role/service-role/securespherecloudtrailrole

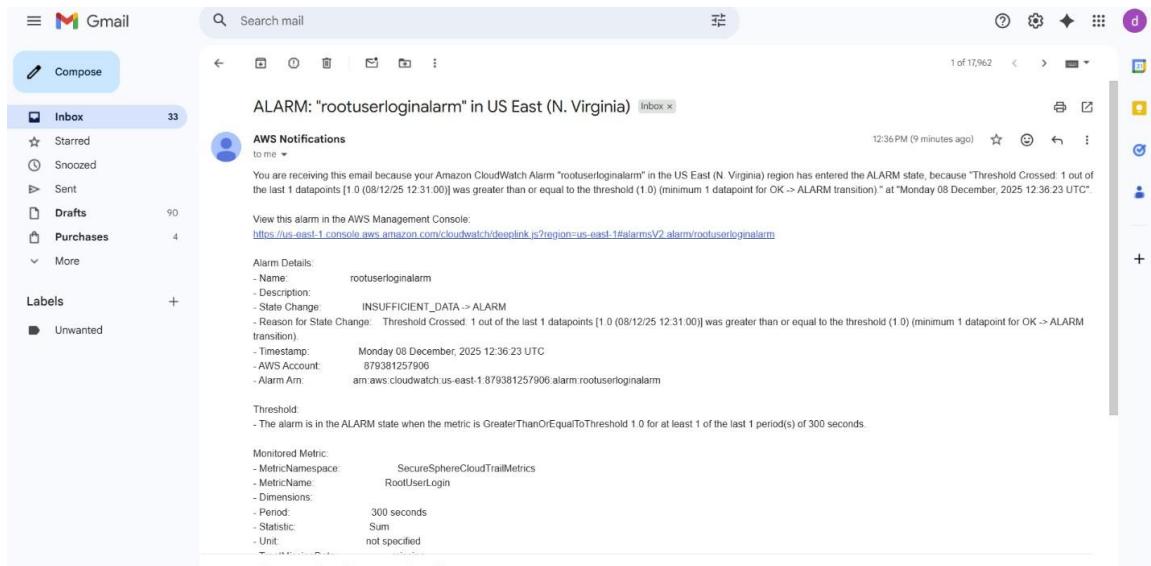
SNS notifications were configured to alert the security team when suspicious events were detected.

7. CloudWatch Metric Filters and Alarms

Metric filters were created for critical security events, including:

- Unauthorized API calls
- Console login without MFA
- Root account usage - IAM policy changes
- CloudTrail configuration changes
- AWS Config changes
- S3 bucket policy changes

Each metric was linked to a CloudWatch alarm and SNS notification.



Manual test events were generated to validate metric activation and alert delivery.

8. Security Hub Integration

AWS Security Hub was enabled to aggregate findings from multiple services, including:

- AWS Config
- GuardDuty
- Inspector
- Macie

Security standards were reviewed and findings were prioritized based on risk level.

9. S3 Security Monitoring and Remediation

A test S3 bucket was created with public access enabled to validate detection capabilities.

GuardDuty and Security Hub detected the misconfiguration.

Block public access settings are disabled for the S3 bucket

Finding: Block public access settings are disabled for the S3 bucket

Severity: High

Details: All bucket-level block public access settings were disabled for the S3 bucket. Access to the bucket is controlled by ...

Resources (1)

Type	Policy IAM User/S3 Block Public Access Disabled
Region	us-east-1
Account	879381257906
Age	3 hours
Created time	December 08, 2025, 16:43 (UTC+00:00)

Public access was removed to remediate the high severity finding.

Overview

Finding trends are an average count of findings for the last 90 days. Filters on the findings table do not affect the graph.

Finding count

Compliance

Status	Pass
Control	S3.8
Config rule	securityhub-s3-bucket-level-public-access-prohibited-f5c4e845
Standards	standards/aws-foundational-security-best-practices/v/1.0.0
Requirements	standards/cis-aws-foundations-benchmark/v/5.0.0
CIS AWS Foundations Benchmark	v5.0.0/2.1.4.2

Remediation

For information on how to correct this issue, consult the AWS Security Hub controls documentation.

Server access logging was enabled to improve audit visibility.

10. Network Security and VPC Flow Logs

VPC Flow Logs were enabled for all production VPCs.

Logs were sent to CloudWatch Log Groups for monitoring and analysis.

An EC2 instance was used to generate network traffic for validation.

Both ACCEPT and REJECT traffic was verified in log streams.

11. Network Monitoring Alerts

Metric filters were created for:

- Rejected traffic events
- Log delivery failures

The screenshot shows the AWS CloudWatch Metrics Filter configuration page. It displays two metric filters: 'RejectLogFilter' and 'LogdeliveryFailure'. The 'RejectLogFilter' has a filter pattern of 'REJECT' and applies to the metric 'SecureSphereVPCFlowLogsMonitoring / RejectedTraffic' with a value of 1. The 'LogdeliveryFailure' filter has a filter pattern of 'NODATA' or 'SKIPDATA' and applies to the metric 'SecureSphereVPCFlowLogs / logFailureDelivery' with a value of 1.

CloudWatch alarms were configured to notify the security team of network anomalies.

The screenshot shows an email from AWS Notifications in a Gmail inbox. The subject is 'ALARM: "FlowLogsalarm" in US East (N. Virginia)'. The email body contains details about the alarm threshold being crossed, the timestamp (Tuesday 09 December, 2025 22:38:08 UTC), and the AWS account information. It also includes a link to the AWS Management Console and a section titled 'Alarm Details' with specific metric details.

12. Documentation and Governance

All configurations, alerting workflows, and remediation procedures were documented.

This documentation supported audits, incident response, and knowledge transfer.

13. Outcomes and Impact

This implementation delivered the following improvements:

- Centralized threat detection - Faster incident response
- Improved network visibility

- Automated remediation
- Reduced security misconfigurations
- Enhanced compliance monitoring

14. Conclusion

This monitoring and threat detection framework provides comprehensive visibility into cloud security risks.

Through centralized logging, automated response, and continuous analysis, the environment supports secure and resilient operations.

ADEDAYO