

Continuous Monitoring and Threat Detection on AWS .

1. Problem Statement.

A SaaS organization required continuous threat detection across its AWS environment. Limited monitoring and delayed alerting increased security risk. A centralized detection framework was required.

I implemented this solution to improve visibility into API activity, network traffic, security misconfigurations, and threat indicators, while enabling rapid incident response.

2. Objectives

The primary objectives were to:

- Enable continuous security monitoring
- Detect suspicious IAM and API activity
- Identify network-based threats
- Centralize security findings
- Automate low-severity remediation
- Improve incident response readiness

3. Security Monitoring Architecture

The monitoring framework was built using the following AWS services:

- Amazon GuardDuty
- AWS CloudTrail
- AWS Security Hub
- VPC Flow Logs
- Amazon CloudWatch
- Amazon SNS
- AWS Lambda

All security events and findings were centralized for analysis and response.

4. GuardDuty Configuration

I enabled Amazon GuardDuty to monitor:

The screenshot shows the AWS GuardDuty Summary page. On the left, a sidebar lists various protection plans: S3 Protection, EKS Protection, Extended Threat Detection (marked as new), Runtime Monitoring, Malware Protection, RDS Protection, and Lambda Protection. Below this are sections for Accounts, Usage, Suppression rules (marked as new), Settings, and Lists. A "What's New" section is also present. The main content area is titled "Summary" and includes an "Overview" section with four boxes: "Attack sequences - new" (0), "Total findings" (0), "Resources with findings" (0), and "Accounts with findings" (0). Below this is a "Findings" section with a grid for Critical, Medium, and Low severity levels, all showing 0 findings. At the bottom of the findings grid are buttons for "Top threats" and "Top attack sequences only". To the right of the findings grid is a promotional box for the AWS Security Hub, which highlights prioritized risk detection, potential attack path visualization, central management of Security Hub CSPM, GuardDuty, and Inspector, and cost optimization through unified pricing. It also features "Compare pricing" and "Try Security Hub" buttons.

- CloudTrail logs
- VPC Flow Logs
- DNS activity

GuardDuty was configured to generate findings for suspicious behavior and compromised credentials.

An SNS topic was created to distribute GuardDuty alerts to the security team.

The screenshot shows the AWS SNS Topics page. The left sidebar lists Dashboard, Topics (selected), Subscriptions, Mobile (Push notifications, Text messaging (SMS)), and a New Feature announcement about High Throughput FIFO topics. The main content area shows a topic named "GuarddutyFindings". The "Details" section includes fields for Name (GuarddutyFindings), ARN (arn:aws:sns:us-east-1:879381257906:GuarddutyFindings), Display name (-), and Type (Standard). Below this are tabs for Subscriptions, Access policy, Data protection policy, Delivery policy (HTTP/S), Delivery status logging, Encryption, and Tags. The "Subscriptions" tab is selected, showing one subscription with ID 1, Endpoint (a Lambda function ARN), Status (Active), and Protocol (Lambda). There are buttons for Edit, Delete, Request confirmation, Confirm subscription, and Create subscription.

Sample findings were generated to validate alert delivery.

5. Automated Remediation for Low-Severity Findings

I implemented automated remediation for low-severity findings using AWS Lambda.

```

Code source Info
EXPLORER LOWSEVERITYGUARDDUTYREMEDIATION lambda_function.py ...
DEPLOY ✓ Current Deploy (Ctrl+Shift+U) Test (Ctrl+Shift+I)
TEST EVENTS [NONE SELECTED] + Create new test event
lambda_function.py
1 import json
2 import boto3
3
4 ec2 = boto3.client('ec2')
5
6 def lambda_handler(event, context):
7     print("Received event: " + json.dumps(event))
8     finding = event['detail']
9     severity = finding['severity']
10
11     # Only process low severity
12     if severity > 4:
13         print("Skipping non low severity finding")
14         return
15
16     resources = finding.get('resources', [])
17     for resource in resources:
18         if resource['type'] == 'AWS::EC2::Instance':
19             instance_id = resource['id']
20             sg_response = ec2.describe_instances(instanceIds=[instance_id])
21             security_groups = sg_response['Reservations'][0]['Instances'][0]['securityGroups']
22
23             ip_address = finding.get('service', {}).get('action', {}).get('remoteipDetails', {}).get('ipAddressV4')
24             if ip_address:
25                 for sg in security_groups:
26                     sg_id = sg['groupId']
27                     try:
28                         ec2.revoke_security_group_ingress(
    
```

A dedicated IAM role was created with permissions to modify tagged security groups.

Step 1: Modify permissions in Lambdalowseverityremediationrole

Step 2: Review and save

Modify permissions in Lambdalowseverityremediationrole Info

Add permissions by selecting services, actions, resources, and conditions. Build permission statements using the JSON editor.

Policy editor

```

1 {
2     "Version": "2012-10-17",
3     "Statement": [
4         {
5             "Effect": "Allow",
6             "Action": [
7                 "ec2:DescribeInstances",
8                 "ec2:DescribeSecurityGroups",
9                 "ec2:RevokeSecurityGroupIngress",
10                "ec2:AuthorizeSecurityGroupIngress",
11                "logs:CreateLogGroup",
12                "logs:CreateLogStream",
13                "logs:PutLogEvents"
14            ],
15            "Resource": "*"
16        }
17    ]
18 }
    
```

Visual | **JSON** Actions ▾

Edit statement Remove

Add actions

Choose a service

Included CloudWatch Logs EC2

Available AI Operations AMP API Gateway API Gateway V2 ARC Region switch ARC Zonal Shift

In production environments, permissions would be restricted to specific ARNs and tagged resources to enforce least privilege.

A Lambda function was deployed and triggered using EventBridge when low-severity findings occurred.

6. CloudTrail and API Monitoring

I configured a centralized CloudTrail trail to capture all AWS API activity.

Logs were forwarded to CloudWatch Logs for near real-time analysis.

The screenshot shows the AWS CloudTrail Trail configuration page for the trail 'SecureSphereCloudtrail'. The 'General details' section includes:

- Trail logging:** Logging is enabled.
- Trail name:** SecureSphereCloudtrail.
- Multi-region trail:** Yes.
- Apply trail to my organization:** Not enabled.

Log file validation: Enabled. Last file validation delivered on December 08, 2025, at 11:20:19 UTC.

SNS notification delivery: Disabled. Last SNS notification was not provided.

CloudWatch Logs: Log group 'aws-cloudtrail-logs-securesphere' is associated with the trail.

IAM Role: The role 'arm:iam::879381257906:role/service-role/securespherecloudtrailrole' is assigned to the trail.

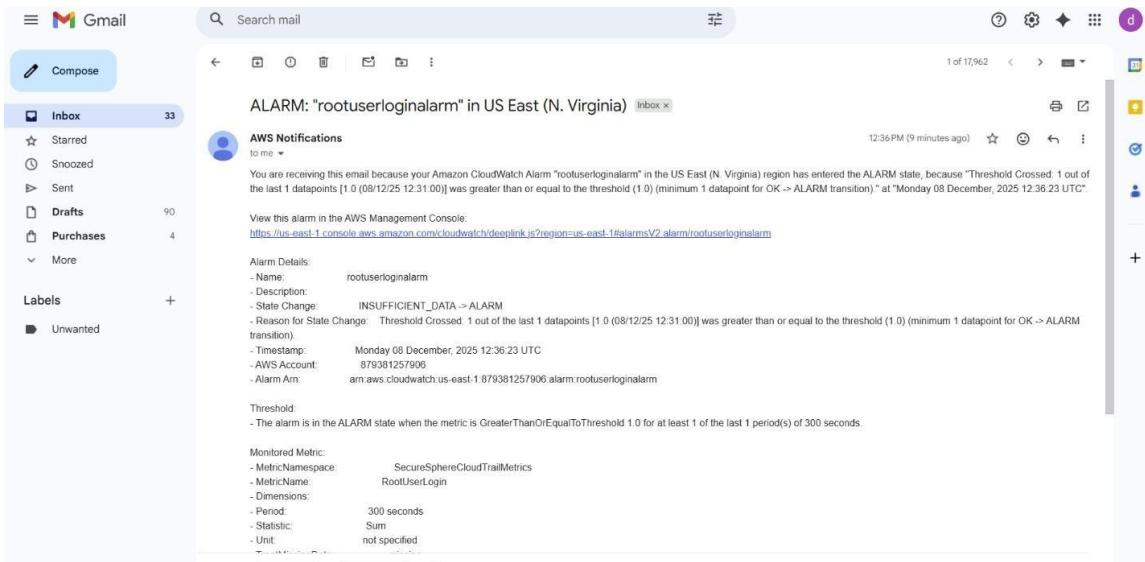
SNS notifications were configured to alert the security team when suspicious events were detected.

7. CloudWatch Metric Filters and Alarms

Metric filters were created for critical security events, including:

- Unauthorized API calls
- Console login without MFA
- Root account usage - IAM policy changes
- CloudTrail configuration changes
- AWS Config changes
- S3 bucket policy changes

Each metric was linked to a CloudWatch alarm and SNS notification.



Manual test events were generated to validate metric activation and alert delivery.

8. Security Hub Integration

AWS Security Hub was enabled to aggregate findings from multiple services, including:

- AWS Config
- GuardDuty
- Inspector
- Macie

Security standards were reviewed and findings were prioritized based on risk level.

9. S3 Security Monitoring and Remediation

A test S3 bucket was created with public access enabled to validate detection capabilities.

GuardDuty and Security Hub detected the misconfiguration.

Block public access settings are disabled for the S3 bucket

Finding: Block public access settings are disabled for the S3 bucket

Severity: High

Details:

- S3 general purpose buckets should block public access
- VPC default security groups should not allow inbound or outbound traffic
- VPC default security groups should not allow inbound or outbound traffic
- A Kubernetes API commonly used in Impact tactics invoked from a Tor exit node IP address.

Resources

Type	Policy: IAMUser/S3BlockPublicAccessDisabled
Region	us-east-1
Account	879381257906
Age	3 hours
Created time	December 06, 2025, 16:43 (UTC+00:00)

Public access was removed to remediate the high severity finding.

Compliance

Status	Pass
Control	S3.8
Config rule	securityhub-s3-bucket-level-public-access-prohibited-f5c4e845
Standards	standards/aws-foundational-security-best-practices/v/1.0.0
Requirements	CIS AWS Foundations Benchmark v5.0.0/2.1.4.2

Remediation

For information on how to correct this issue, consult the AWS Security Hub controls documentation.

Server access logging was enabled to improve audit visibility.

10. Network Security and VPC Flow Logs

VPC Flow Logs were enabled for all production VPCs.

The screenshot shows the AWS VPC dashboard under 'Your VPCs'. A green success message at the top says 'Successfully created flow log for the following resource: vpc-0d0014fc2012e7b67'. Below it, a card for a flow log named 'fl-0e52ae20a963c6135' is displayed. The card includes details like 'Destination Type: cloud-watch-logs', 'Destination Name: SecureSphereVpcFlowLogs', 'Traffic Type: All', 'Max Aggregation Interval: 1 minute', and 'Log Format: Default'. It also shows the 'Creation Time' as 'Tuesday 9 December 2025 at 21:36:55 GMT'. The 'Tags' and 'Integrations' tabs are visible below the main card.

Logs were sent to CloudWatch Log Groups for monitoring and analysis.

An EC2 instance was used to generate network traffic for validation.

Both ACCEPT and REJECT traffic was verified in log streams.

The screenshot shows the AWS CloudWatch Log management interface. A green success message at the top says 'Log group "SecureSphereVpcFlowLogs" has been created.' Below it, a table titled 'Log events' lists several log entries. The columns are 'Timestamp' and 'Message'. The messages show network traffic details, such as IP addresses and port numbers, followed by 'REJECT OK' or 'ACCEPT OK'. The log entries are timestamped from December 9, 2025, at 09:21:47 to 09:21:48.

11. Network Monitoring Alerts

Metric filters were created for:

- Rejected traffic events
- Log delivery failures

The screenshot shows the AWS CloudWatch Metrics Filter interface. At the top, a green banner indicates that a metric filter named "LogdeliveryFailure" has been created. The interface includes tabs for Log streams, Tags, Anomaly detection, Metric filters (which is selected), Subscription filters, Contributor Insights, Data protection, and Field in. Below the tabs, there's a search bar and a list of filters. Two filters are currently displayed: "RejectLogFilter" and "LogdeliveryFailure".

- RejectLogFilter:**
 - Filter pattern: "REJECT"
 - Field selection criteria: Metric: SecureSphereVPCFlowLogsMonitoring / RejectedTraffic
 - Metric value: 1
 - Default value: -
 - Applied on transformed logs: -
 - Unit: -
- LogdeliveryFailure:**
 - Filter pattern: "NODATA" "SKIPDATA"
 - Field selection criteria: Metric: SecureSphereVPCFlowLogs / logFailedDelivery
 - Metric value: 1
 - Default value: -
 - Applied on transformed logs: -
 - Unit: -

CloudWatch alarms were configured to notify the security team of network anomalies.

The screenshot shows a Gmail inbox with 44 unread messages. A specific email from "AWS Notifications" is highlighted, titled "ALARM: "FlowLogsalarm" in US East (N. Virginia)". The email content details an alarm triggered by the "FlowLogsalarm" metric in the US East (N. Virginia) region. It provides information about the threshold crossed (1 out of 1 datapoint), the time of the alarm (Tuesday 09 December, 2025 22:38:08 UTC), and the alarmArn (arn:aws:cloudwatch:us-east-1:879381257906:alarm:FlowLogsalarm). It also lists the monitored metric details: MetricNamespace: SecureSphereVPCFlowLogsMonitoring, MetricName: RejectedTraffic, Dimensions: Period: 300 seconds, Statistic: Sum.

12. Documentation and Governance

All configurations, alerting workflows, and remediation procedures were documented.

This documentation supported audits, incident response, and knowledge transfer.

13. Outcomes and Impact

This implementation delivered the following improvements:

- Centralized threat detection - Faster incident response
- Improved network visibility

- Automated remediation
- Reduced security misconfigurations
- Enhanced compliance monitoring

14. Conclusion

This monitoring and threat detection framework provides comprehensive visibility into cloud security risks.

Through centralized logging, automated response, and continuous analysis, the environment supports secure and resilient operations.

ADDEDAYO